

Carnegie Mellon University Africa

04-800K: AIOps: Continuous and Automated IT and AI Monitoring

Assynath Thompson Mlay, [amlaytho](#)

MSIT 2025

Lab 5: Capturing application metrics with Istio service mesh.

Lab Preparations:

```
● asnath@ssynath3r:~/AIops_Labs/Lab 5$ ./provision2.sh boutique amlaytho us-west1

No current clusters found, continuing to deploy one.
Default change: VPC-native is the default mode during cluster creation for versions greater than 1.21.0-gke.1500. To create advanced routes based clusters, please pass the '--no-enable-ip-alias' flag.
Default change: During creation of nodepools or autoscaling configuration changes for cluster versions greater than 1.24.1-gke.800 a default location policy is applied. For Spot and PMM it defaults to ANY, and for all other VM kinds a BALANCED policy is used. To change the default values use the '--location-policy' flag.
Note: Your Pod address range ('--cluster-ipv4-cidr') can accommodate at most 1008 node(s).
Creating cluster boutique in us-west1-a... Cluster is being health-checked (Kubernetes Control Plane is healthy)...done.
Created [https://container.googleapis.com/v1/projects/amlaytho/zones/us-west1-a/clusters/boutique].
To inspect the contents of your cluster, go to: https://console.cloud.google.com/kubernetes/workload/_gccloud/us-west1-a/boutique?project=amlaytho
Kubeconfig entry generated for boutique.
NAME      LOCATION  MASTER_VERSION  MASTER_IP      MACHINE_TYPE  NODE_VERSION    NUM_NODES  STATUS
boutique  us-west1-a  1.30.5-gke.1014001  34.105.25.193  e2-standard-4  1.30.5-gke.1014001  1          RUNNING



✓ Istio core installed 🚢
✓ Istiod installed 🍷
✓ Egress gateways installed 🌐
✓ Ingress gateways installed 🌐
✓ Installation complete
cluster-wide operations.
istiod 1/1 1 35s
istio install complete
```

Made this installation the default for

Figure 1: Provisioning and downloading Istio

```
serviceaccount/adservice created
serviceaccount/cartservice created
serviceaccount/checkoutservic created
serviceaccount/currencyservice created
serviceaccount/emailservice created
serviceaccount/frontend created
serviceaccount/loadgenerator created
serviceaccount/paymentservice created
serviceaccount/productcatalogservice created
serviceaccount/recommendationservice created
serviceaccount/shippingservice created
service/adservice created
service/cartservice created
service/checkoutservic created
service/currencyservice created
service/emailservice created
service/frontend created
service/paymentservice created
service/productcatalogservice created
service/recommendationservice created
service/redis-cart created
service/shippingservice created
deployment.apps/adservice created
deployment.apps/cartservice created
deployment.apps/checkoutservic created
deployment.apps/currencyservice created
deployment.apps/emailservice created
deployment.apps/frontend created
deployment.apps/loadgenerator created
deployment.apps/paymentservice created
deployment.apps/productcatalogservice created
deployment.apps/recommendationservice created
deployment.apps/redis-cart created
deployment.apps/shippingservice created
gateway.gateway.networking.k8s.io/istio-gateway created
httproute.gateway.networking.k8s.io/frontend-route created
serviceentry.networking.istio.io/allow-egress-google-metadata created
serviceentry.networking.istio.io/allow-egress-googleapis created
```

Figure 2: Microservices deployed

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
adservice	ClusterIP	34.118.238.182	<none>	9555/TCP	22s
cartservice	ClusterIP	34.118.226.158	<none>	7070/TCP	21s
checkoutservice	ClusterIP	34.118.235.212	<none>	5050/TCP	21s
currencyservice	ClusterIP	34.118.232.8	<none>	7000/TCP	20s
emailservice	ClusterIP	34.118.232.12	<none>	5000/TCP	19s
frontend	ClusterIP	34.118.227.242	<none>	80/TCP	19s
istio-gateway-istio	LoadBalancer	34.118.226.185	<pending>	15021:30154/TCP,80:31048/TCP	6s
kubernetes	ClusterIP	34.118.224.1	<none>	443/TCP	3m43s
paymentservice	ClusterIP	34.118.231.154	<none>	50051/TCP	18s
productcatalogservice	ClusterIP	34.118.239.4	<none>	3550/TCP	17s
recommendationservice	ClusterIP	34.118.236.117	<none>	8080/TCP	16s
redis-cart	ClusterIP	34.118.239.61	<none>	6379/TCP	16s
shippingservice	ClusterIP	34.118.225.40	<none>	50051/TCP	15s

serviceaccount/grafana created
 configmap/grafana created
 service/grafana created
 deployment.apps/grafana created
 configmap/istio-grafana-dashboards created
 configmap/istio-services-grafana-dashboards created
 deployment.apps/jaeger created
 service/tracing created
 service/zipkin created
 service/jaeger-collector created
 serviceaccount/kiali created
 configmap/kiali created
 clusterrole.rbac.authorization.k8s.io/kiali created
 clusterrolebinding.rbac.authorization.k8s.io/kiali created
 role.rbac.authorization.k8s.io/kiali-controlplane created
 rolebinding.rbac.authorization.k8s.io/kiali-controlplane created
 service/kiali created
 deployment.apps/kiali created
 serviceaccount/loki created
 configmap/loki created
 configmap/loki-runtime created
 service/loki-memberlist created
 service/loki-headless created
 service/loki created
 statefulset.apps/loki created

Figure 3: Clusters deployed with services

rolebinding.rbac.authorization.k8s.io/kiali-controlplane created					
service/kiali created					
deployment.apps/kiali created					
serviceaccount/loki created					
configmap/loki created					
configmap/loki-runtime created					
service/loki-memberlist created					
service/loki-headless created					
service/loki created					
statefulset.apps/loki created					
serviceaccount/prometheus created					
configmap/prometheus created					
clusterrole.rbac.authorization.k8s.io/prometheus created					
clusterrolebinding.rbac.authorization.k8s.io/prometheus created					
service/prometheus created					
deployment.apps/prometheus created					
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
grafana	ClusterIP	34.118.234.87	<none>	3000/TCP	33s
istio-egressgateway	ClusterIP	34.118.239.106	<none>	80/TCP,443/TCP	104s
istio-ingressgateway	LoadBalancer	34.118.234.132	34.169.104.12	15021:30714/TCP,80:31598/TCP,443:31500/TCP,31400:32486/TCP,15443:31003/TCP	104s
istiod	ClusterIP	34.118.227.209	<none>	15010/TCP,15012/TCP,443/TCP,15014/TCP	114s
jaeger-collector	ClusterIP	34.118.226.141	<none>	14268/TCP,14250/TCP,9411/TCP,4317/TCP,4318/TCP	24s
kiali	ClusterIP	34.118.225.86	<none>	20001/TCP,9090/TCP	17s
loki	ClusterIP	34.118.228.127	<none>	3100/TCP,9095/TCP	9s
loki-headless	ClusterIP	None	<none>	3100/TCP	10s
loki-memberlist	ClusterIP	None	<none>	7946/TCP	11s
prometheus	ClusterIP	34.118.234.1	<none>	9090/TCP	3s
tracing	ClusterIP	34.118.229.234	<none>	80/TCP,16685/TCP	26s
zipkin	ClusterIP	34.118.226.212	<none>	9411/TCP	25s

Figure 4: Other clusters

Lab Task1: Install Prometheus and Grafana for the Boutique Application

```
asnath@AssynathJr:~/AIOps_Labs/Lab 5$ istioctl dashboard prometheus  
http://localhost:9090
```

Figure 5: Opening Prometheus Dashboard

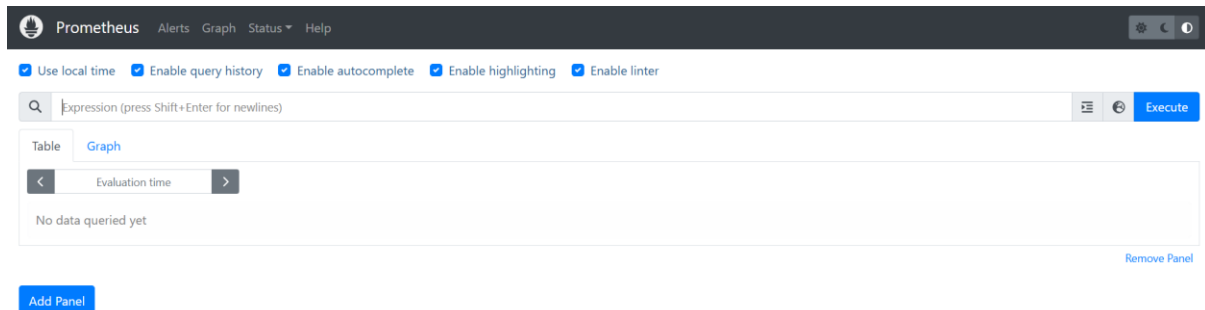


Figure 6: Opened Prometheus

```
asnath@AssynathJr:~/AIOps_Labs/Lab 5$ istioctl dashboard grafana  
http://localhost:3000
```

Figure 7: Opening Grafana Dashboard

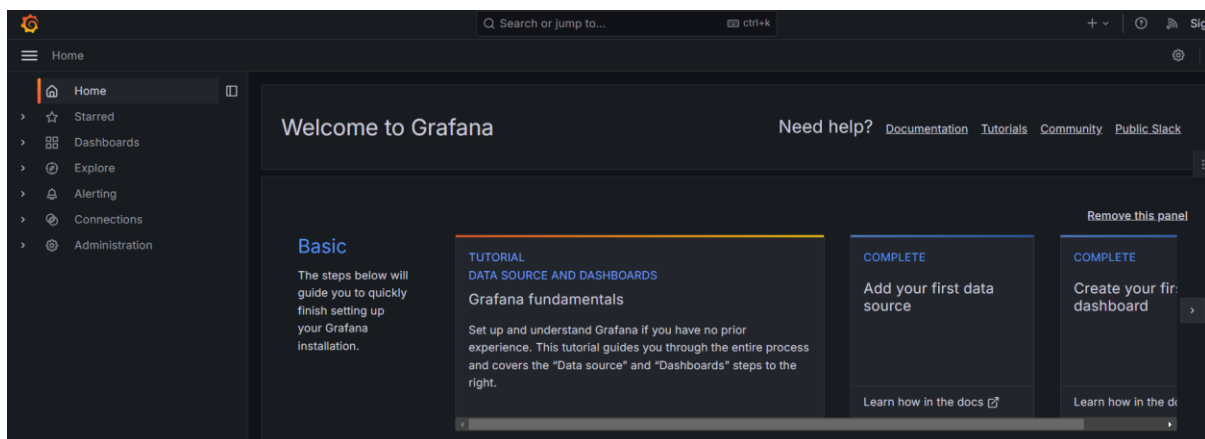


Figure 8: Opening Grafana

```
asnath@AssynathJr:~/AIOps_Labs/Lab 5$ istioctl dashboard kiali  
http://localhost:20001/kiali
```

Figure 9: Opening Kiali Dashboard

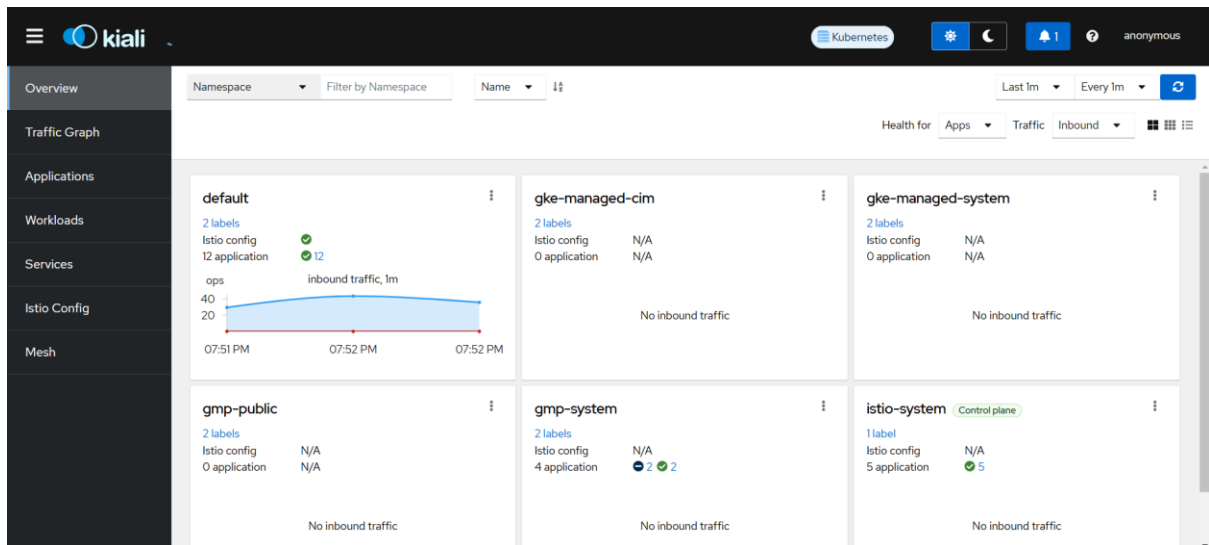


Figure 10: Kiali Dashboard

Running the “istio_requests_total”:

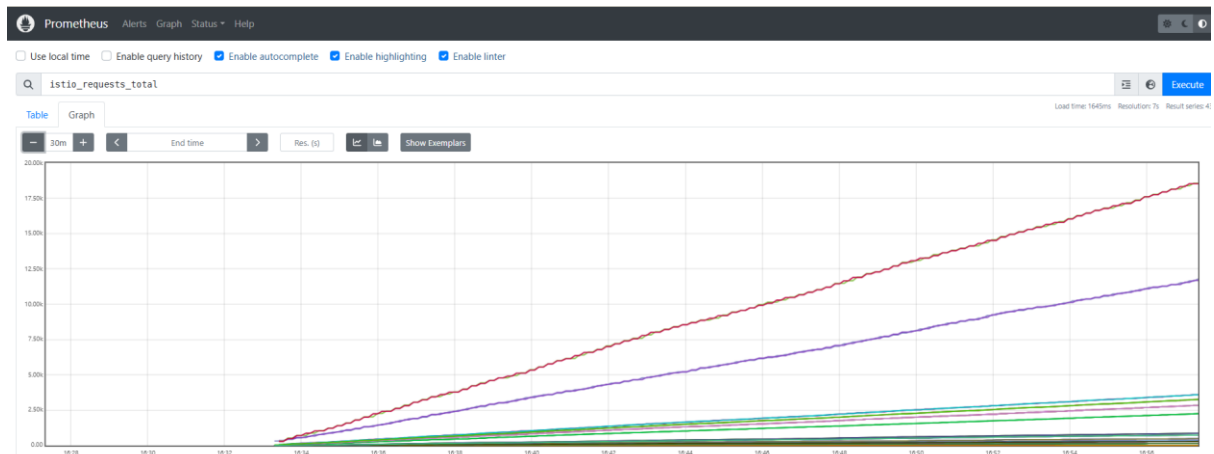


Figure 11: istio_requests_total on Prometheus

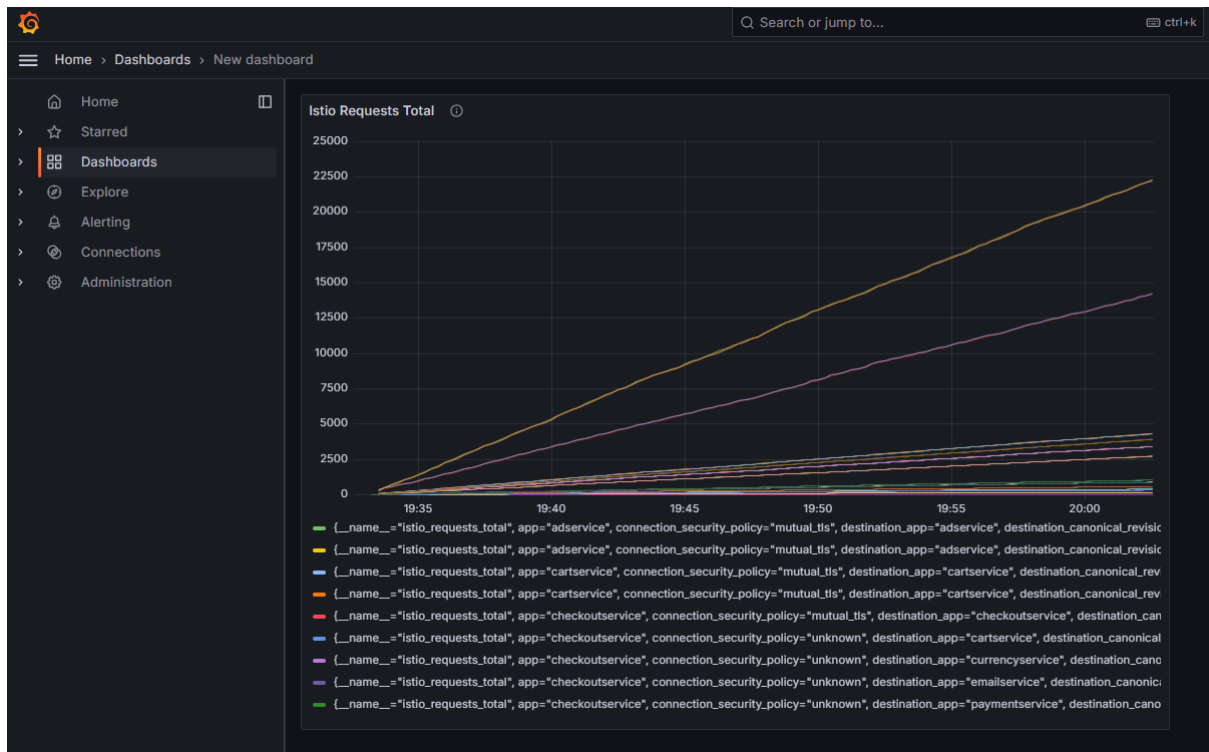


Figure 12: Istio Requests Total on Grafana

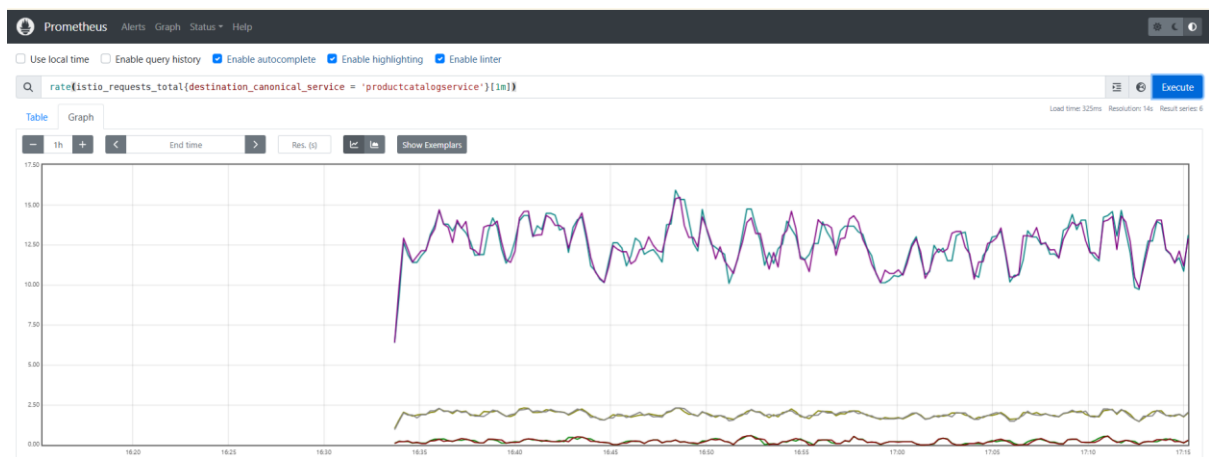


Figure 13: identifying source/destination services



Figure 14: finding the destination services on grafana

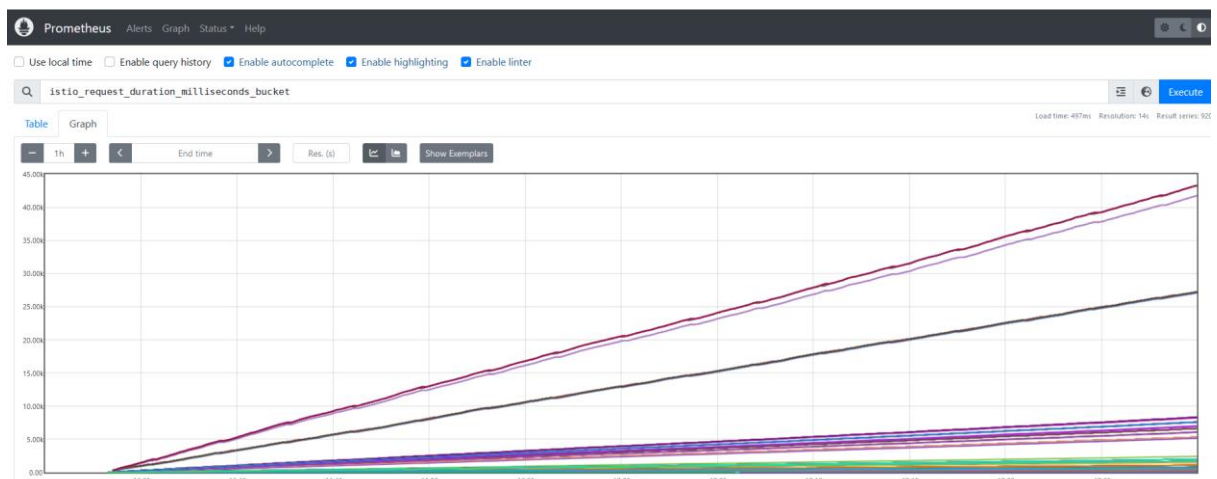


Figure 15: istio request duration milliseconds bucket

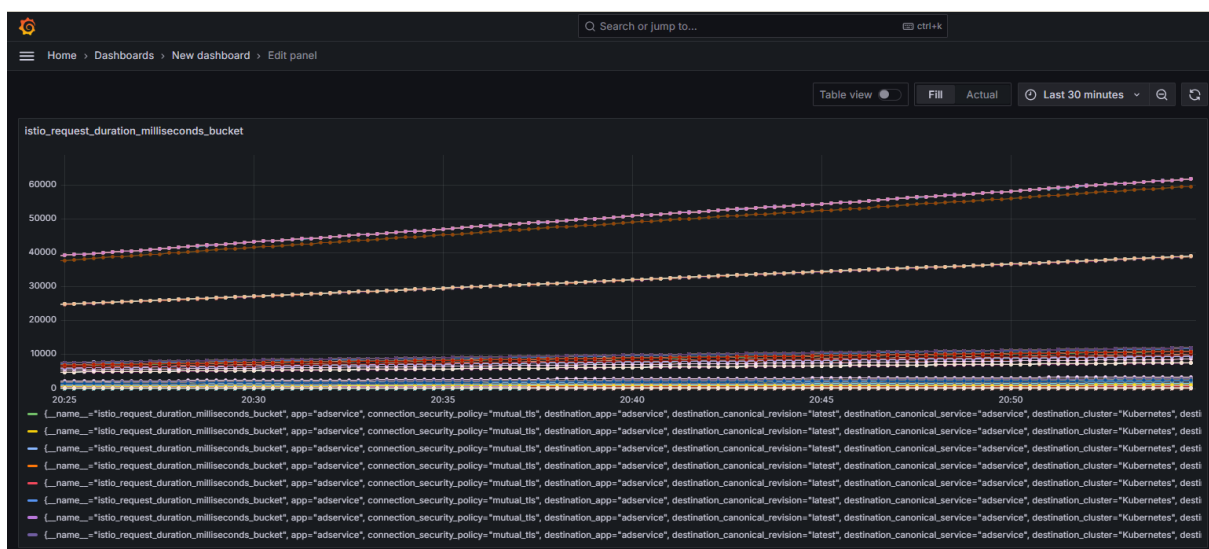


Figure 16: istio request duration milliseconds bucket

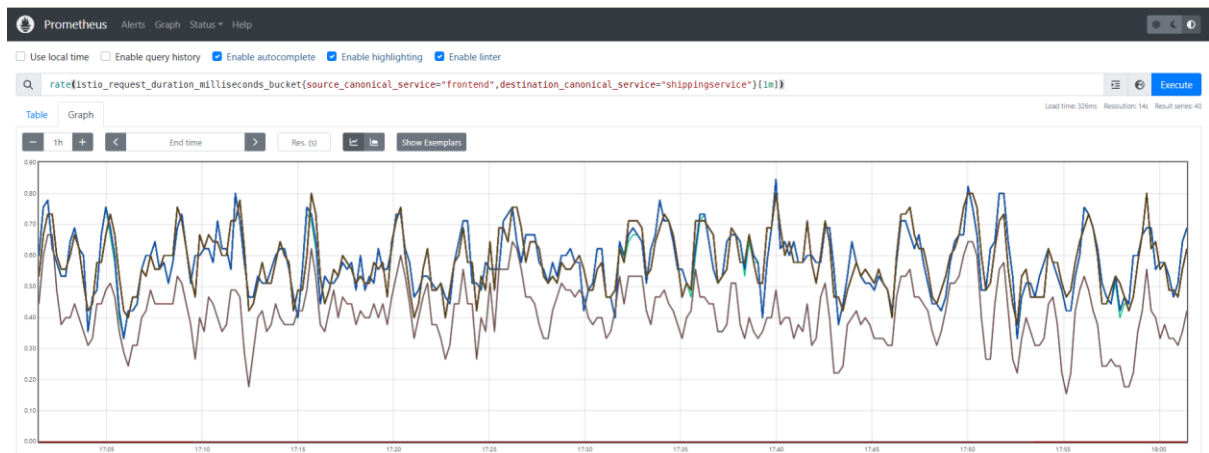


Figure 17: both source and destination canonical



Figure 18: both source and destination canonicals on Grafana

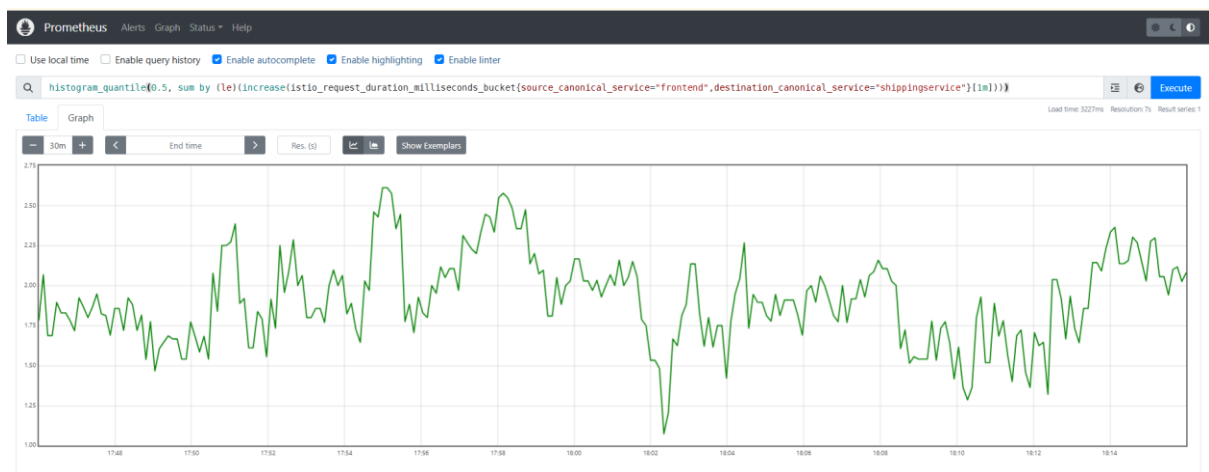


Figure 19: Histogram for the 0.5 percentile on Prometheus


```

● asnath@AssynathJr:~/AIOps_Labs/Lab 5$ kubectl get deployments
NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
adservice                          1/1      1              1            113m
cartservice                        1/1      1              1            113m
checkoutservice                    1/1      1              1            113m
currencyservice                    1/1      1              1            113m
emailservice                       1/1      1              1            113m
frontend                           1/1      1              1            113m
istio-gateway-istio                1/1      1              1            113m
loadgenerator                       1/1      1              1            113m
paymentservice                     1/1      1              1            113m
productcatalogservice              1/1      1              1            113m
recommendationservice              1/1      1              1            113m
redis-cart                         1/1      1              1            113m
shippingservice                    1/1      1              1            113m
○ asnath@AssynathJr:~/AIOps_Labs/Lab 5$ █

```

Figure 22: Checking the deployments.

```

● asnath@AssynathJr:~/AIOps_Labs/Lab 5$ kubectl scale deployment loadgenerator --replicas=2
deployment.apps/loadgenerator scaled
● asnath@AssynathJr:~/AIOps_Labs/Lab 5$ kubectl get deployments
NAME                                READY    UP-TO-DATE    AVAILABLE    AGE
adservice                          1/1      1              1            115m
cartservice                        1/1      1              1            115m
checkoutservice                    1/1      1              1            115m
currencyservice                    1/1      1              1            115m
emailservice                       1/1      1              1            115m
frontend                           1/1      1              1            115m
istio-gateway-istio                1/1      1              1            115m
loadgenerator                       2/2      2              2            115m
paymentservice                     1/1      1              1            115m
productcatalogservice              1/1      1              1            115m
recommendationservice              1/1      1              1            115m
redis-cart                         1/1      1              1            115m
shippingservice                    1/1      1              1            115m

```

Figure 23: Load generator scaled.

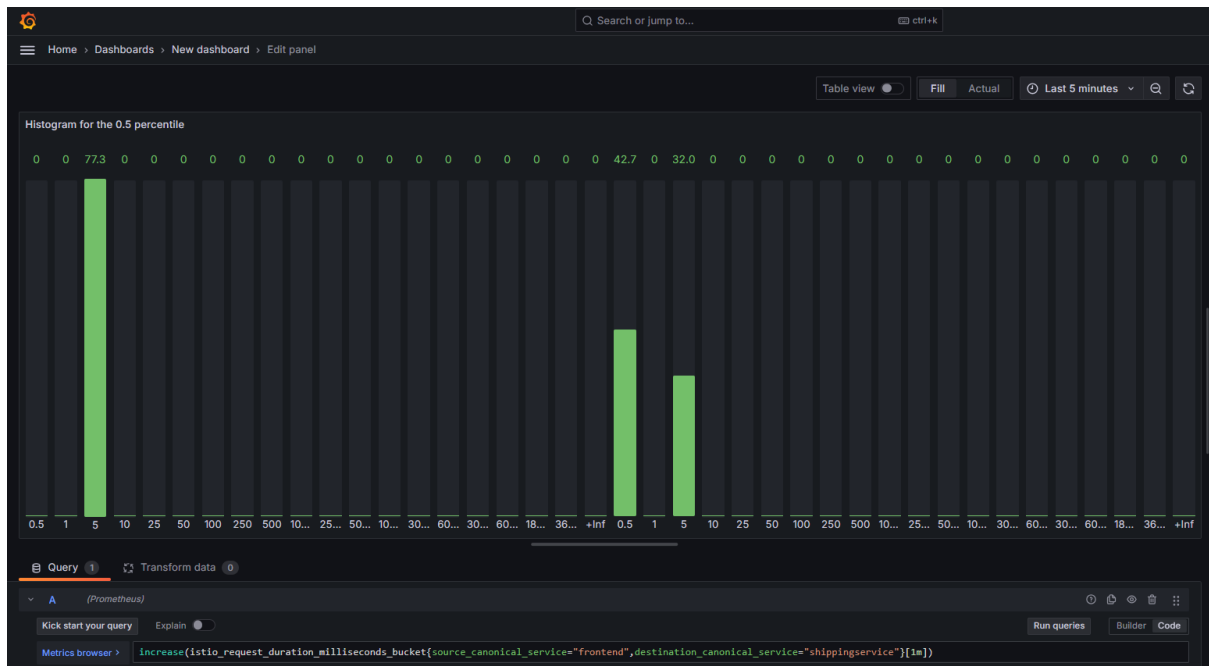
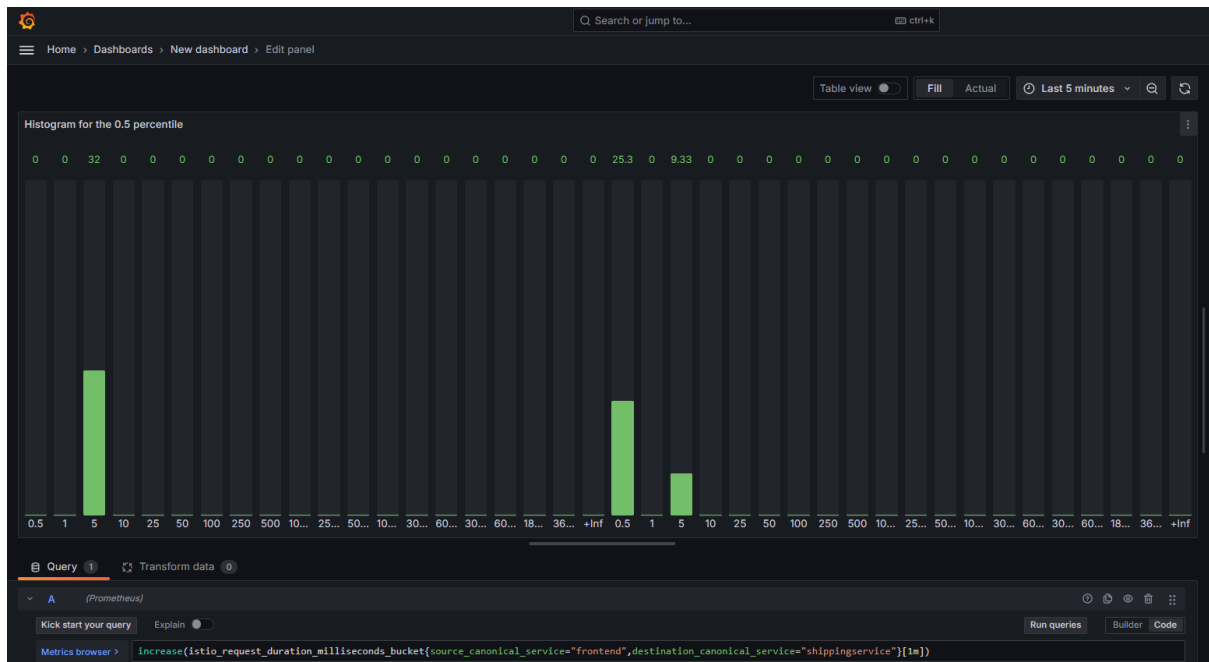


Figure 24: After scaling

```
● asnath@AssynathJr:~/AIOps_Labs/Lab 5$ kubectl scale deployment loadgenerator --replicas=1
deployment.apps/loadgenerator scaled
```

```
asnath@AssynathJr:~/AIOps_Labs/Lab 5$ kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
adservice	1/1	1	1	122m
cartservice	1/1	1	1	122m
checkoutservice	1/1	1	1	121m
currencyservice	1/1	1	1	121m
emailservice	1/1	1	1	121m
frontend	1/1	1	1	121m
istio-gateway-istio	1/1	1	1	121m
loadgenerator	1/1	1	1	121m
paymentservice	1/1	1	1	121m
productcatalogservice	1/1	1	1	121m
recommendationservice	1/1	1	1	121m
redis-cart	1/1	1	1	121m
shippingservice	1/1	1	1	121m



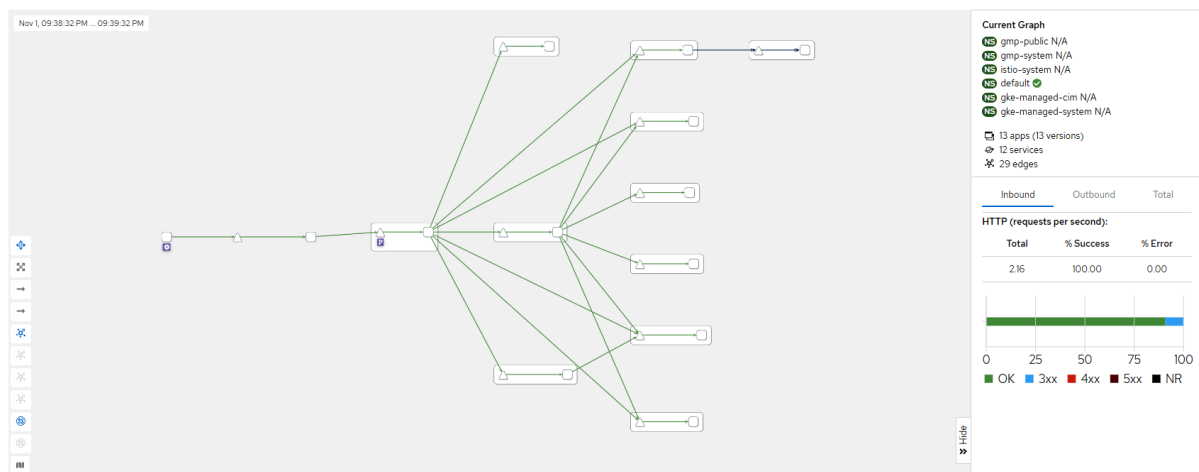
Why do duplicated entries appear for requests between services?

- Duplicates happen because Prometheus tracks each request using multiple labels, like the names and versions of both the sending and receiving services. Each unique combination of these labels creates a new data entry, even if it's the same request. To reduce these duplicates, we can use a function like `sum by (le)` to combine them, or we can filter by just one label, such as the source or destination, for a single view.

Convert to Incremental Rate (1m) and Reduce Duplicates:

- To make cumulative data show changes per minute, we use `rate(metric_name[1m])`, which calculates the request rate over each 1 minute. We can focus on a specific label that limits data to one source-to-destination pair to avoid duplicate entries, giving us a clearer view without extra data.
-

Lab Task 2: Install the Kiali cluster visualization dashboard.



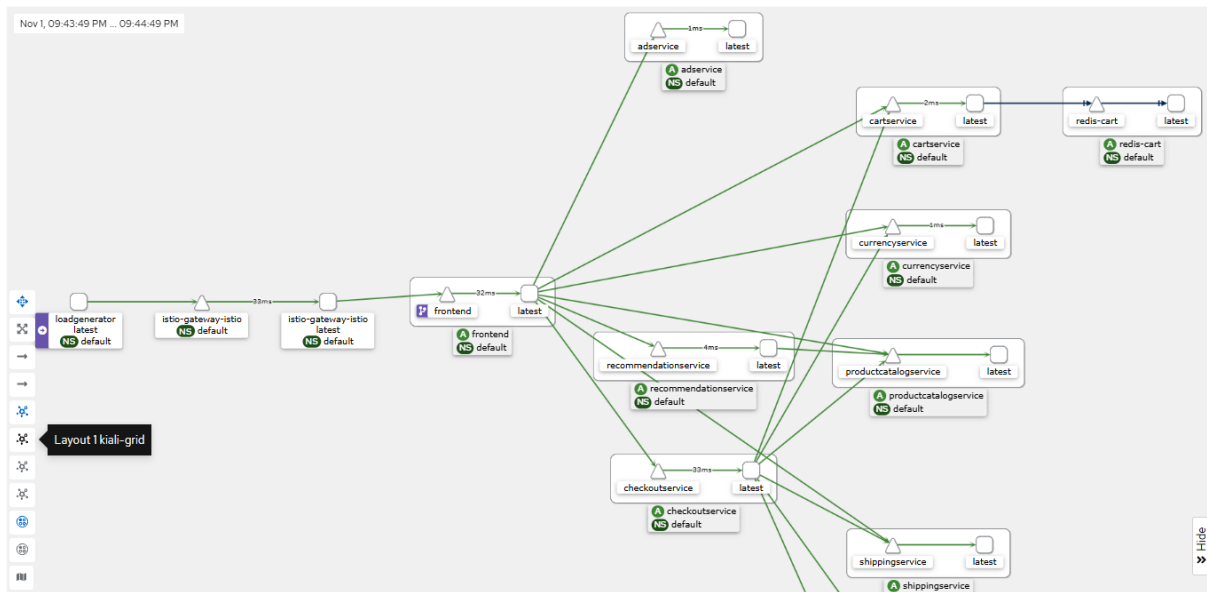


Figure 28: Zoomed in Visualization

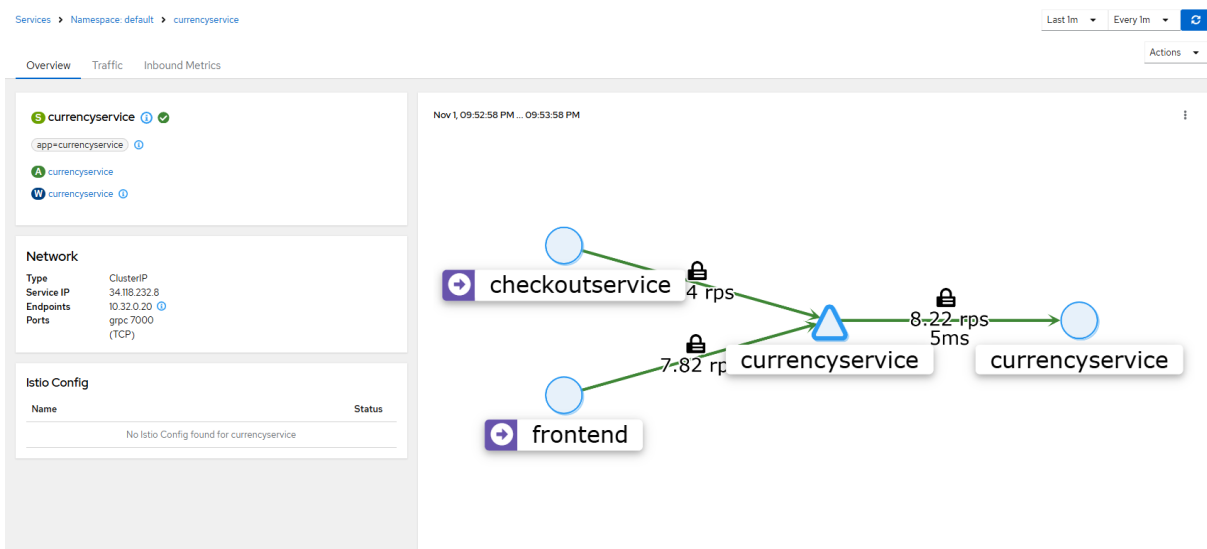


Figure 29: Currency Service

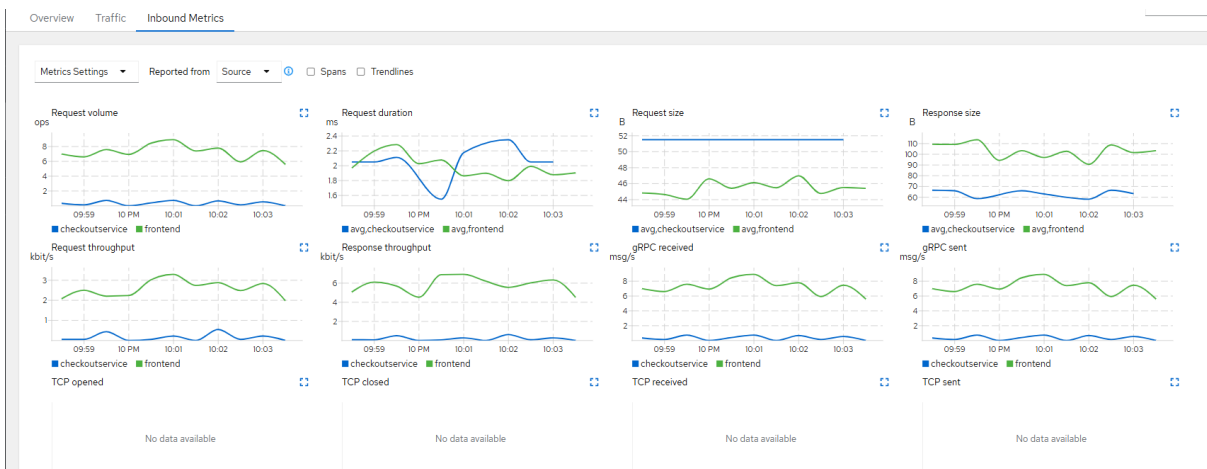


Figure 30: inbound metrics for currency service

Impact of Fault Injection Observed in Kiali:

- After adding a fault injection, we can see delays in Kiali by looking at the Inbound Metrics for `shippingservice`. Requests coming from `frontend` show longer response times, while requests from other services, like `checkoutservice`, don't have delays. This shows that the fault injection only affected requests from `frontend`, as intended.

Lab Task 3: Inject selective delay faults to the Boutique shipping service.

```
asnth@AssynthJr:~/AI0ps_Labs/Lab 5$ kubectl describe deployment frontend
Name: frontend
Namespace: default
CreationTimestamp: Fri, 01 Nov 2024 19:32:04 +0300
Labels: app=frontend
Annotations: deployment.kubernetes.io/revision: 1
Selector: app=frontend
Replicas: 1 desired | 1 updated | 1 total | 1 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels: app=frontend
  Annotations: sidecar.istio.io/rewriteAppHTTPProbers: true
  Service Account: frontend
  Containers:
    server:
      Image: us-central1-docker.pkg.dev/google-samples/microservices-demo/frontend:v0.10.1
      Port: 8080/TCP
      Host Port: 0/TCP
      Limits:
        cpu: 200m
        memory: 128Mi
      Requests:
        cpu: 100m
        memory: 64Mi
      Liveness: http-get http://:8080/_healthz delay=10s timeout=1s period=10s #success=1 #failure=3
      Readiness: http-get http://:8080/_healthz delay=10s timeout=1s period=10s #success=1 #failure=3
      Environment:
        PORT: 8080
        PRODUCT_CATALOG_SERVICE_ADDR: productcatalogservice:3550
        CURRENCY_SERVICE_ADDR: currencyservice:7000
        CART_SERVICE_ADDR: cartservice:7070
        RECOMMENDATION_SERVICE_ADDR: recommendationservice:8080
        SHIPPING_SERVICE_ADDR: shippingservice:50051
        CHECKOUT_SERVICE_ADDR: checkoutservice:5050
        AD_SERVICE_ADDR: adservice:9555
        SHOPPING_ASSISTANT_SERVICE_ADDR: shoppingassistantservice:80
        ENABLE_PROFILER: 0
      Mounts: <none>
      Volumes: <none>
      Node-Selectors: <none>
      Tolerations: <none>
  Conditions:
    Type          Status  Reason
    ----          -
    Available      True    MinimumReplicasAvailable
    Progressing    True    NewReplicaSetAvailable
    OldReplicaSets: <none>
    NewReplicaSet:  frontend-cb9967686 (1/1 replicas created)
    Events:        <none>
```

Figure 31: Kubectl getting the description of the frontend.

```

● asnath@AssynathJr:~/AIops_Labs/Lab 5$ kubectl apply -f lab5_faultinjection.yaml
virtualservice.networking.istio.io/shippingservice created

```

```

○ asnath@AssynathJr:~/AIops_Labs/Lab 5$

```

Figure 32: Creating a yaml for fault injection

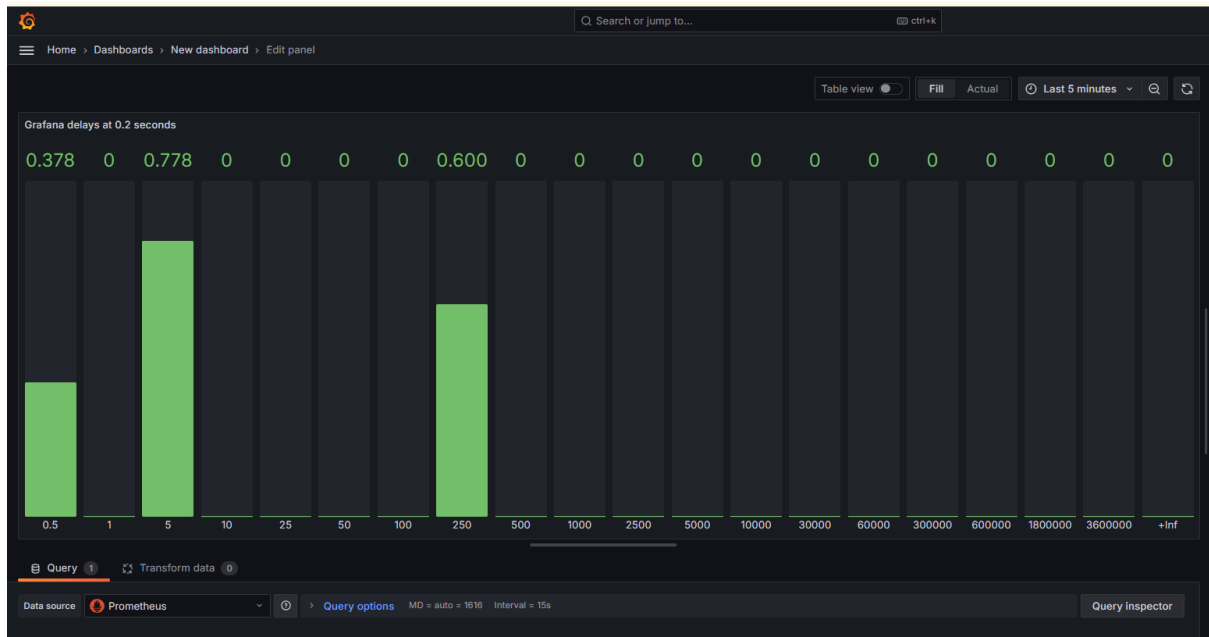


Figure 33: Grafana delays by 0.2 seconds

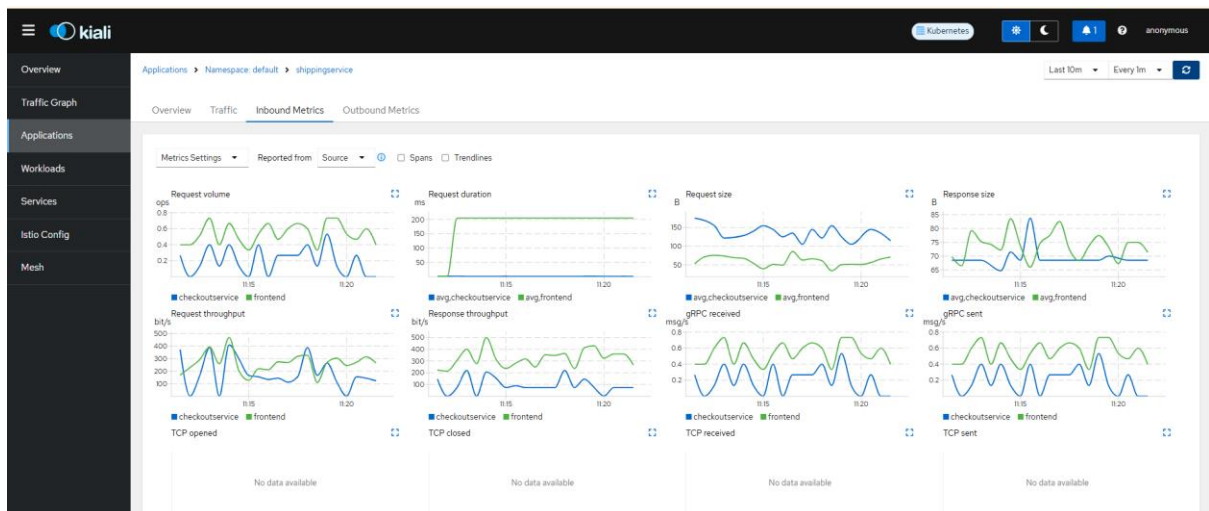


Figure 34: Kiali delays

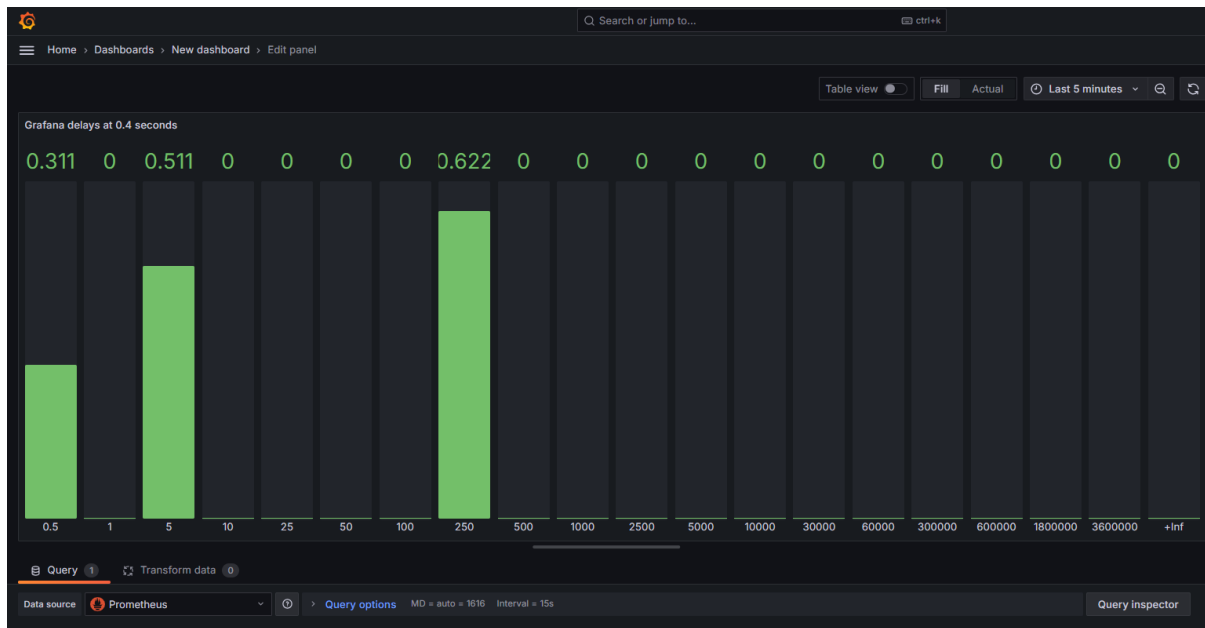


Figure 35: Grafana delays at 0.4

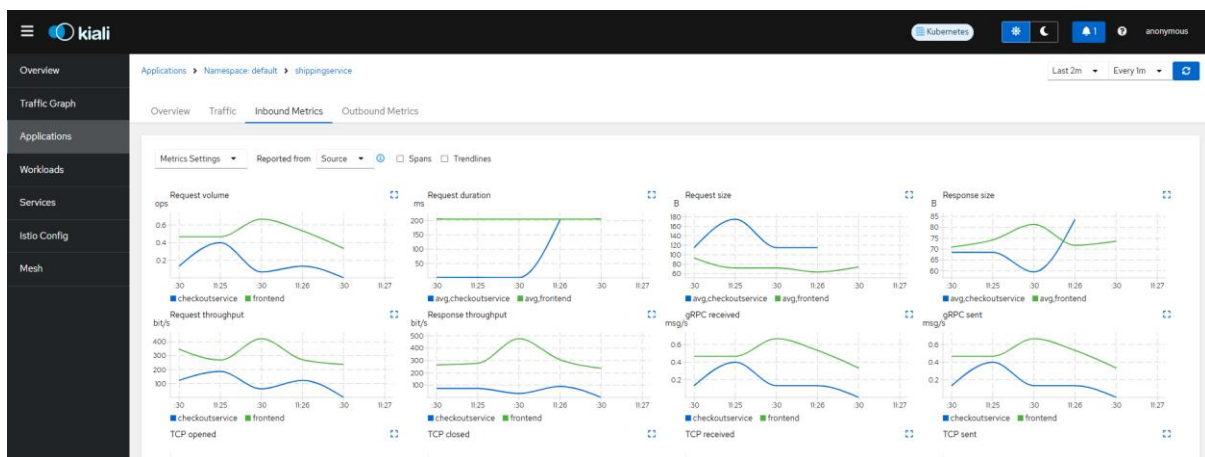


Figure 36: 0.4 delays on Kiali. Where the request duration jumped to meet the frontend service.

Grafana Histogram with Both Delays:

- In Grafana, the histogram for `shippingservice` shows two peaks: one at 0.2 seconds for delays from `frontend` and another at 0.4 seconds for delays from `checkoutservice` on specific requests like "GetQuote." Using `sum by (le)` groups these delay data points into one histogram, so we don't see duplicates from Istio's added delay labels.

Where is the Delay Imposed (Kiali)?

- In Kiali, looking at the Source metrics for `shippingservice` shows that the delay affects the whole path from `frontend` to `shippingservice`. This means that Istio applies the delay at the network level before it reaches the application code, as we can see by comparing the Source and Destination metrics.